

Mlflow screen shots

Mlflow screenshot showing experiment with all features to determine top model considerations.

The screenshot displays the Mlflow web interface. The top navigation bar includes the Mlflow logo (1.26.1), 'Experiments', and 'Models' tabs. A search bar and a 'Share' button are also present. The main content area shows the details for an experiment named 'Experiment_with_all_features'. A sidebar on the left lists experiments: 'Experiment_with_all_f...', 'Experiment_with_red...', and 'Lead_Scoring_Trainin...'. The experiment details section includes a description, a table of runs, and a search bar. The table shows 34 matching runs, with columns for Start Time, Duration, Run Name, User, Source, Models, and various metrics (AUC, Accuracy, F1, Kappa, MCC, Prec., Recall, TT, C). The runs are sorted by Start Time, showing the most recent runs at the top.

Experiment ID: 1

Description [Edit](#)

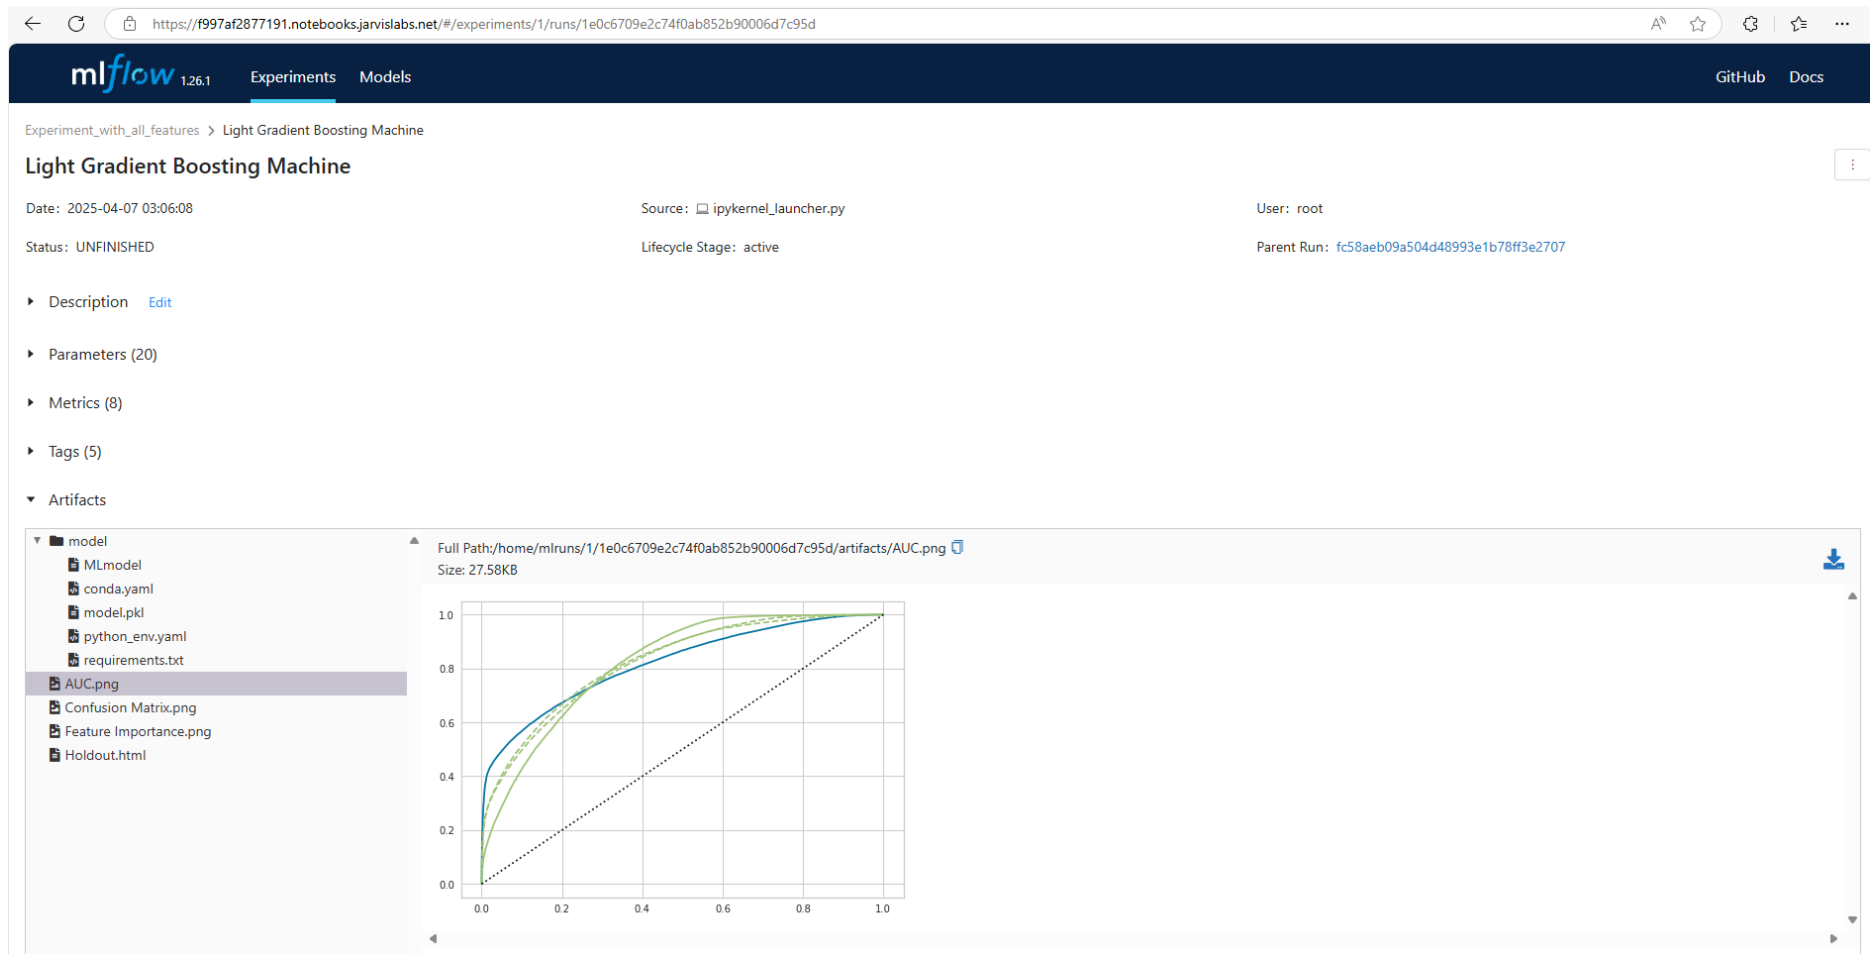
[Refresh](#) [Compare](#) [Delete](#) [Download CSV](#) [Start Time](#) [All time](#)

[Columns](#) Only show differences [Search](#) [Filter](#) [Clear](#)

Showing 34 matching runs

| | | Duration | Run Name | User | Source | Models | Metrics | | | | | | | | Parameters |
|--------------------------|---------------|----------|------------------|------|--------------|---------|---------|----------|-------|-------|-------|-------|--------|-------|------------|
| | Start Time | | | | | | AUC | Accuracy | F1 | Kappa | MCC | Prec. | Recall | TT | C |
| <input type="checkbox"/> | 4 minutes ago | | Session Initi... | root | ipykernel... | - | - | - | - | - | - | - | - | - | - |
| <input type="checkbox"/> | 3 minutes ago | | Light Gradie... | root | ipykernel... | sklearn | 0.821 | 0.738 | 0.762 | 0.476 | 0.484 | 0.703 | 0.831 | 1.37 | - |
| <input type="checkbox"/> | 3 minutes ago | | Ridge Classif... | root | ipykernel... | sklearn | 0 | 0.701 | 0.728 | 0.401 | 0.408 | 0.67 | 0.797 | 0.09 | - |
| <input type="checkbox"/> | 3 minutes ago | | Naive Bayes | root | ipykernel... | sklearn | 0.734 | 0.657 | 0.725 | 0.313 | 0.357 | 0.608 | 0.897 | 0.092 | - |
| <input type="checkbox"/> | 3 minutes ago | | Linear Discr... | root | ipykernel... | sklearn | 0.773 | 0.701 | 0.728 | 0.401 | 0.408 | 0.67 | 0.797 | 0.29 | - |
| <input type="checkbox"/> | 3 minutes ago | | Logistic Reg... | root | ipykernel... | sklearn | 0.784 | 0.71 | 0.74 | 0.419 | 0.43 | 0.673 | 0.823 | 1.674 | 1.0 |
| <input type="checkbox"/> | 3 minutes ago | | Decision Tre... | root | ipykernel... | sklearn | 0.817 | 0.736 | 0.758 | 0.472 | 0.479 | 0.704 | 0.82 | 0.182 | - |
| <input type="checkbox"/> | 3 minutes ago | | Extra Trees C... | root | ipykernel... | sklearn | 0.818 | 0.737 | 0.758 | 0.472 | 0.479 | 0.704 | 0.82 | 1.092 | - |
| <input type="checkbox"/> | 3 minutes ago | | Random For... | root | ipykernel... | sklearn | 0.819 | 0.737 | 0.76 | 0.474 | 0.482 | 0.704 | 0.825 | 0.886 | - |
| <input type="checkbox"/> | 3 minutes ago | | Extreme Gra... | root | ipykernel... | sklearn | 0.821 | 0.738 | 0.761 | 0.476 | 0.483 | 0.703 | 0.829 | 3.314 | - |
| <input type="checkbox"/> | 3 minutes ago | | Light Gradie... | root | ipykernel... | sklearn | 0.821 | 0.738 | 0.762 | 0.476 | 0.484 | 0.703 | 0.831 | 0.562 | - |

Mlflow screen of all artifacts for top performing model (LGBM) when using all features.



Mlflow screenshot showing experiments with selected features

The screenshot shows the Mlflow web interface. The top navigation bar includes the Mlflow logo (1.26.1), 'Experiments', and 'Models' tabs. On the right, there are links for 'GitHub' and 'Docs'. The main header shows the experiment name 'Experiment_with_reduced_features' with a '+ <' icon and a 'Share' button.

Below the header, there is a search bar for 'Search Experiments'. A list of experiments is shown on the left sidebar:

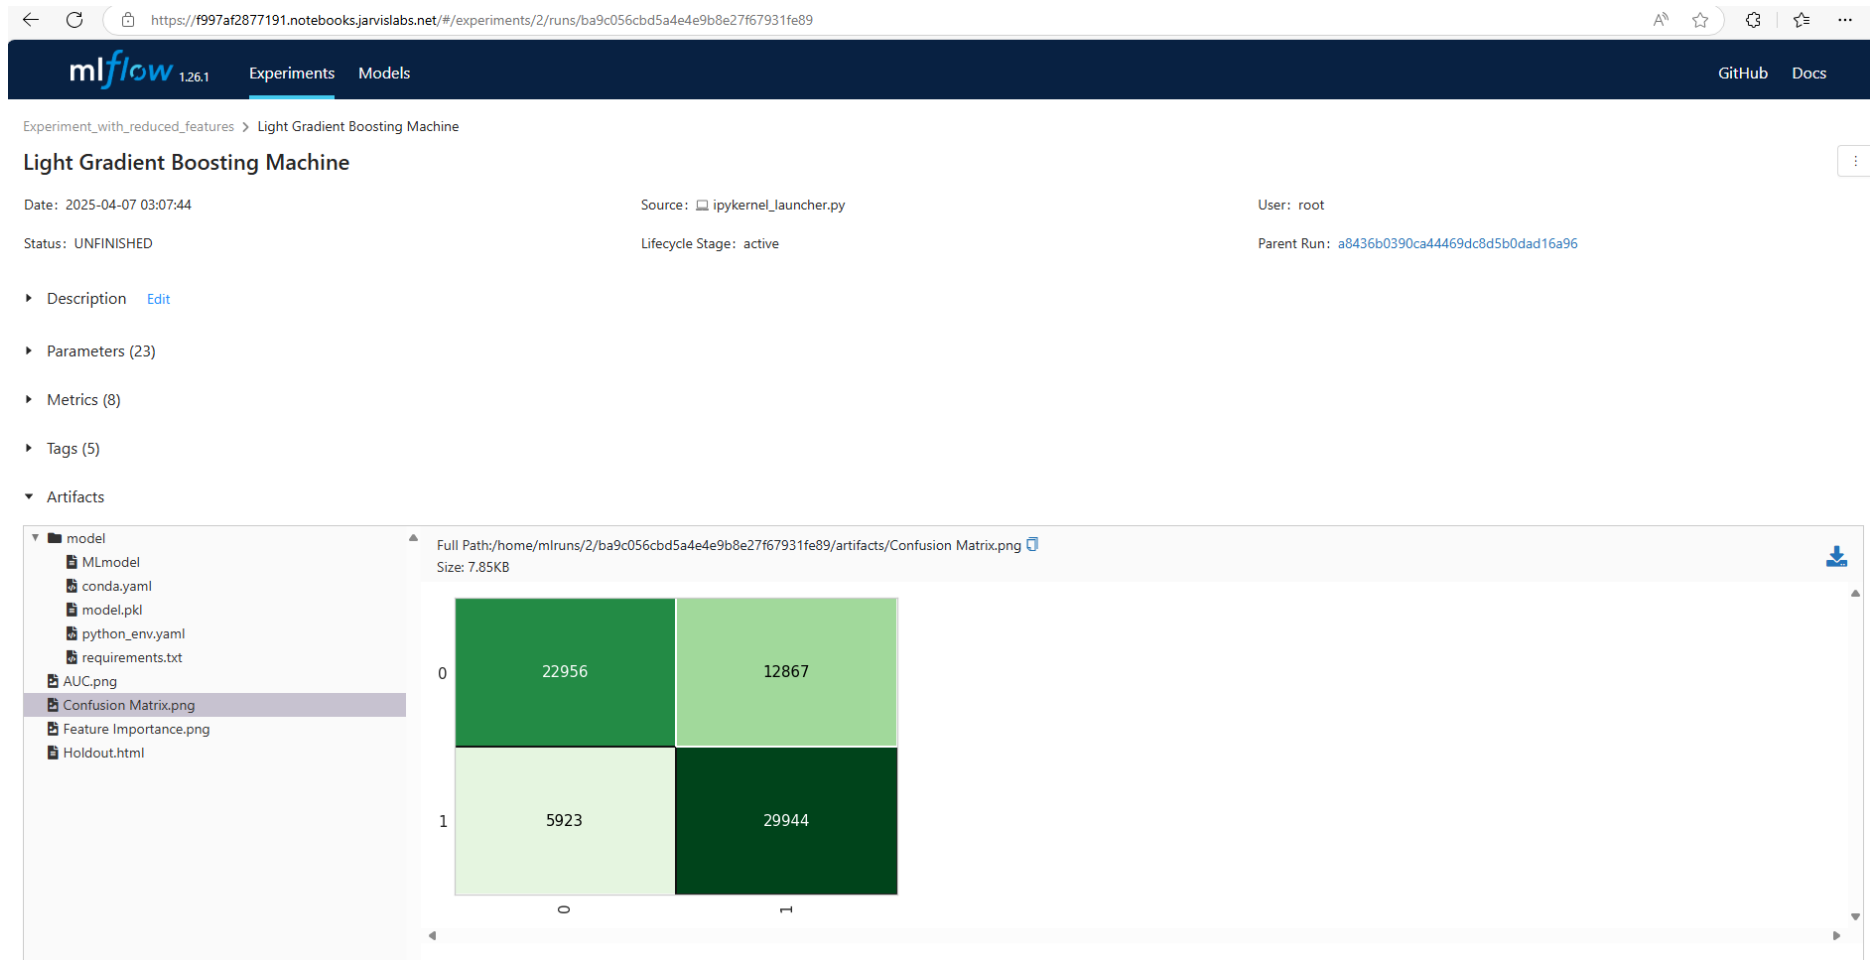
- Experiment_with_all_f... (with a trash icon)
- Experiment_with_red... (selected, with a trash icon)
- Lead_Scoring_Trainin... (with a trash icon)

The main content area displays the experiment details for 'Experiment_with_reduced_features' (Experiment ID: 2). It includes a description section with a 'Description' tab and an 'Edit' link. Below this are action buttons: 'Refresh', 'Compare', 'Delete', 'Download CSV', and a 'Start Time' dropdown set to 'All time'. There is also a 'Columns' button, a toggle for 'Only show differences', and a search bar containing the query 'metrics.rmse < 1 and params.model = "tree"'. The results show 'Showing 23 matching runs'.

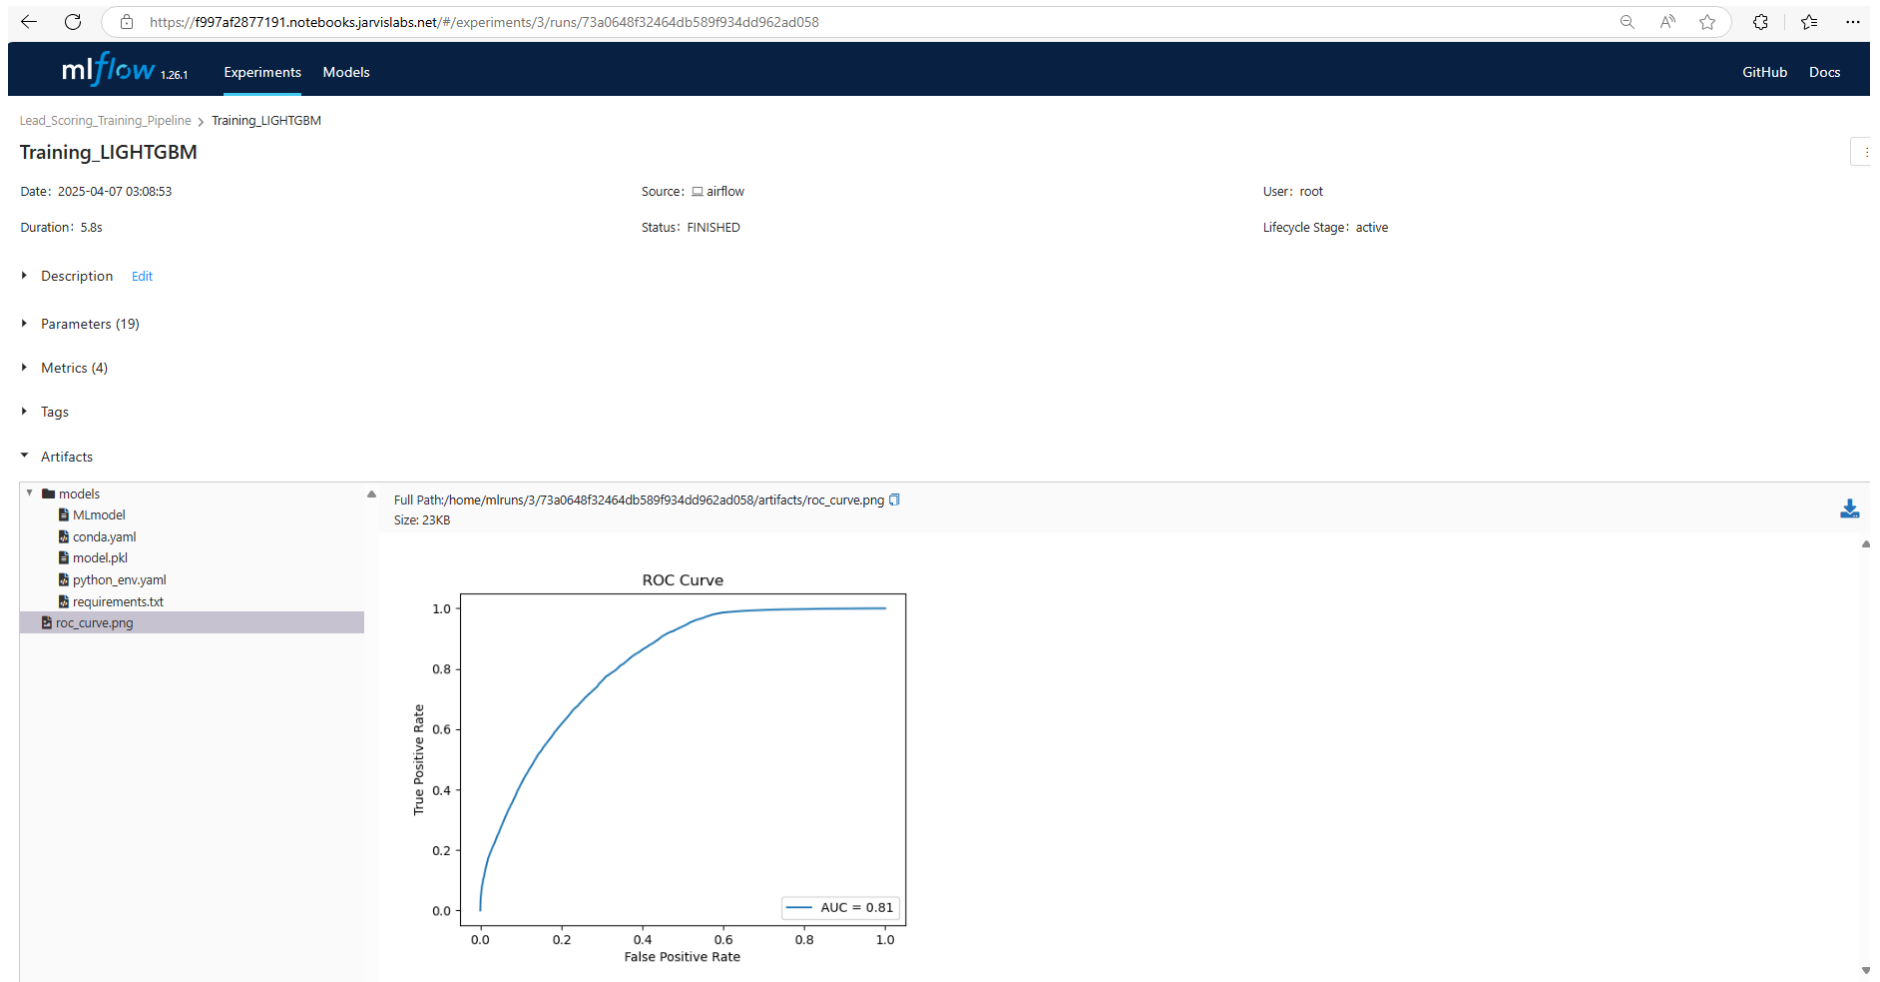
The table below lists the runs, sorted by 'Start Time' (descending). The columns are: Run Name, User, Source, Version, Models, and Metrics (AUC, Accuracy, F1, Kappa, MCC, Prec., Recall).

| | Run Name | User | Source | Version | Models | AUC | Accuracy | F1 | Kappa | MCC | Prec. | Recall |
|--------------------------|------------------|------|--------------|---------|---------|-------|----------|-------|-------|-------|-------|--------|
| <input type="checkbox"/> | Session Initi... | root | ipykernel... | - | - | - | - | - | - | - | - | - |
| <input type="checkbox"/> | Light Gradie... | root | ipykernel... | - | sklearn | 0.82 | 0.738 | 0.762 | 0.476 | 0.485 | 0.703 | 0.832 |
| <input type="checkbox"/> | Light Gradie... | root | ipykernel... | - | sklearn | 0.821 | 0.738 | 0.762 | 0.476 | 0.485 | 0.702 | 0.833 |
| <input type="checkbox"/> | Decision Tre... | root | ipykernel... | - | sklearn | 0.817 | 0.736 | 0.758 | 0.472 | 0.479 | 0.704 | 0.82 |
| <input type="checkbox"/> | Extra Trees C... | root | ipykernel... | - | sklearn | 0.818 | 0.736 | 0.758 | 0.472 | 0.479 | 0.704 | 0.82 |
| <input type="checkbox"/> | Random For... | root | ipykernel... | - | sklearn | 0.819 | 0.737 | 0.759 | 0.473 | 0.48 | 0.704 | 0.823 |
| <input type="checkbox"/> | Extreme Gra... | root | ipykernel... | - | sklearn | 0.821 | 0.738 | 0.761 | 0.475 | 0.484 | 0.702 | 0.831 |
| <input type="checkbox"/> | Light Gradie... | root | ipykernel... | - | sklearn | 0.821 | 0.738 | 0.762 | 0.476 | 0.485 | 0.702 | 0.833 |

Mlflow screenshot showing artifacts for best performing model with selected features



Mlflow screenshot showing model training logged by training pipeline



Mlflow screenshot showing trained model being promoted to “production”

The screenshot displays the MLflow web interface in a browser. The address bar shows the URL: `https://f997af2877191.notebooks.jarvislabs.net/#/models/LightGBM/versions/24`. The MLflow logo and version 1.26.1 are in the top left, with navigation tabs for Experiments and Models. The right side of the header has links for GitHub and Docs.

Breadcrumbs indicate the path: Registered Models > LightGBM > Version 24.

Version 24

Registered At: 2025-04-07 03:08:58 Stage: Production Last Modified: 2025-04-07 03:08:59

Source Run: [Training_LIGHTGBM](#)

▸ Description [Edit](#)

▸ Tags

▼ Schema

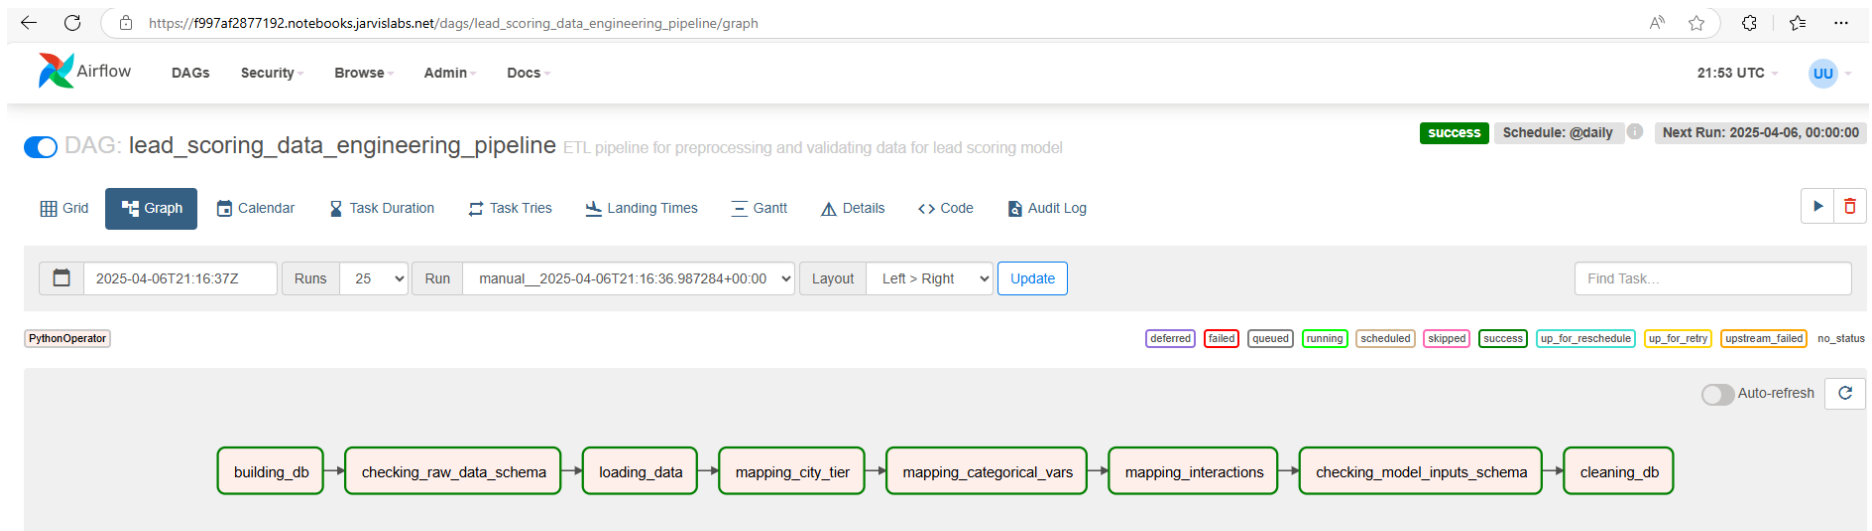
| Name | Type |
|--|------|
| No schema. See MLflow docs for how to include input and output schema with your model. | |

Airflow screenshots

Airflow screenshot showing all DAGs


[illegible]


Airflow screenshot showing successful data pipeline in graph view





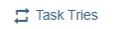
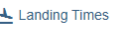








Airflow screenshot showing successful data pipeline in grid view

← ↻ 🔒 https://f997af2877192.notebooks.jarvislabs.net/dags/lead_scoring_data_engineering_pipeline/grid?root= 🔊 ☆ ⚙️ ⚙️ ⋮

 **DAGs** Security Browse Admin Docs 21:53 UTC UU

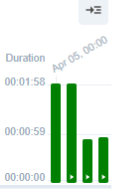
 **DAG: lead_scoring_data_engineering_pipeline** ETL pipeline for preprocessing and validating data for lead scoring model Schedule: @daily ⓘ Next Run: 2025-04-06, 00:00:00









 **Grid**  Graph  Calendar  Task Duration  Task Tries  Landing Times  Gantt  Details  Code  Audit Log  

04/06/2025 09:53:45 PM 📅 25 ▾ All Run Types ▾ All Run States ▾ [Clear Filters](#)

deferred failed queued running scheduled skipped success up_for_reschedule up_for_retry upstream_failed no_status

Auto-refresh ☐


 **Duration** 00:01:58 00:00:50 00:00:00 Apr 05, 00:00

| | |
|------------------------------|---|
| building_db |  |
| checking_raw_data_schema |  |
| loading_data |  |
| mapping_city_tier |  |
| mapping_categorical_vars |  |
| mapping_interactions |  |
| checking_model_inputs_schema |  |
| cleaning_db |  |

DAG
lead_scoring_data_engineering_pipeline

DAG Details

DAG Runs Summary

| | |
|---|--------------------------|
| Total Runs Displayed | 4 |
|  Total success | 4 |
| First Run Start | 2025-04-06, 20:34:13 UTC |
| Last Run Start | 2025-04-06, 21:16:37 UTC |
| Max Run Duration | 00:01:58 |
| Mean Run Duration | 00:01:25 |
| Min Run Duration | 00:00:52 |

DAG Summary

| | |
|-----------------|---|
| Total Tasks | 8 |
| PythonOperators | 8 |

Airflow screenshot showing successful training pipeline in graph view

The screenshot displays the Airflow web interface for a DAG named 'lead_scoring_training_pipeline'. The interface is in 'Graph' view, showing a linear sequence of three tasks: 'start_mlflow_server', 'encoding_categorical_variables', and 'training_model'. All tasks are in a 'success' state, indicated by green borders. The top navigation bar includes links for DAGs, Security, Browse, Admin, and Docs. The top right shows the time as 22:02 UTC and a user icon. The DAG status bar at the top right indicates 'success', a schedule of '@daily', and a next run time of '2025-04-06, 00:00:00'. Below the DAG title, there are tabs for Grid, Graph, Calendar, Task Duration, Task Tries, Landing Times, Gantt, Details, Code, and Audit Log. A filter bar at the top of the graph view shows a date range from '2025-04-06T21:38:30Z' to '2025-04-06T21:38:30Z', with 25 runs. A 'Find Task...' search bar is also present. A legend at the bottom right lists various task statuses: deferred, failed, queued, running, scheduled, skipped, success, up_for_reschedule, up_for_retry, upstream_failed, and no_status. An 'Auto-refresh' toggle is located at the bottom right of the graph area.

https://f997af2877192.notebooks.jarvislabs.net/dags/lead_scoring_training_pipeline/graph

Airflow DAGs Security Browse Admin Docs 22:02 UTC UU

DAG: lead_scoring_training_pipeline Training pipeline for Lead Scoring System success Schedule: @daily Next Run: 2025-04-06, 00:00:00

Grid Graph Calendar Task Duration Task Tries Landing Times Gantt Details Code Audit Log

2025-04-06T21:38:30Z Runs 25 Run manual__2025-04-06T21:38:29.480786+00:00 Layout Left > Right Update Find Task...



PythonOperator deferred failed queued running scheduled skipped success up_for_reschedule up_for_retry upstream_failed no_status


Auto-refresh













```
graph LR; start_mlflow_server --> encoding_categorical_variables; encoding_categorical_variables --> training_model;
```

Airflow screenshot showing successful training pipeline in grid view

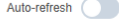

← ↻ 📄 https://f997af2877192.notebooks.jarvislabs.net/dags/lead_scoring_training_pipeline/grid?root= 🔊 ☆ ⚙️ ⚙️ ⋮

 Airflow **DAGs** Security ▾ Browse ▾ Admin ▾ Docs ▾ 22:03 UTC 

 **DAG: lead_scoring_training_pipeline** Training pipeline for Lead Scoring System Schedule: @daily ⓘ Next Run: 2025-04-06, 00:00:00


 Grid  Graph  Calendar  Task Duration  Task Tries  Landing Times  Gantt  Details  <> Code  Audit Log  

04/06/2025 10:03:39 PM 📅 25 ▾ All Run Types ▾ All Run States ▾ [Clear Filters](#)

 Auto-refresh 

Duration
00:00:46
00:00:23
00:00:00
Apr 06, 20:36


start_milflow_server
encoding_categorical_variables
training_model



DAG
lead_scoring_training_pipeline

DAG Details

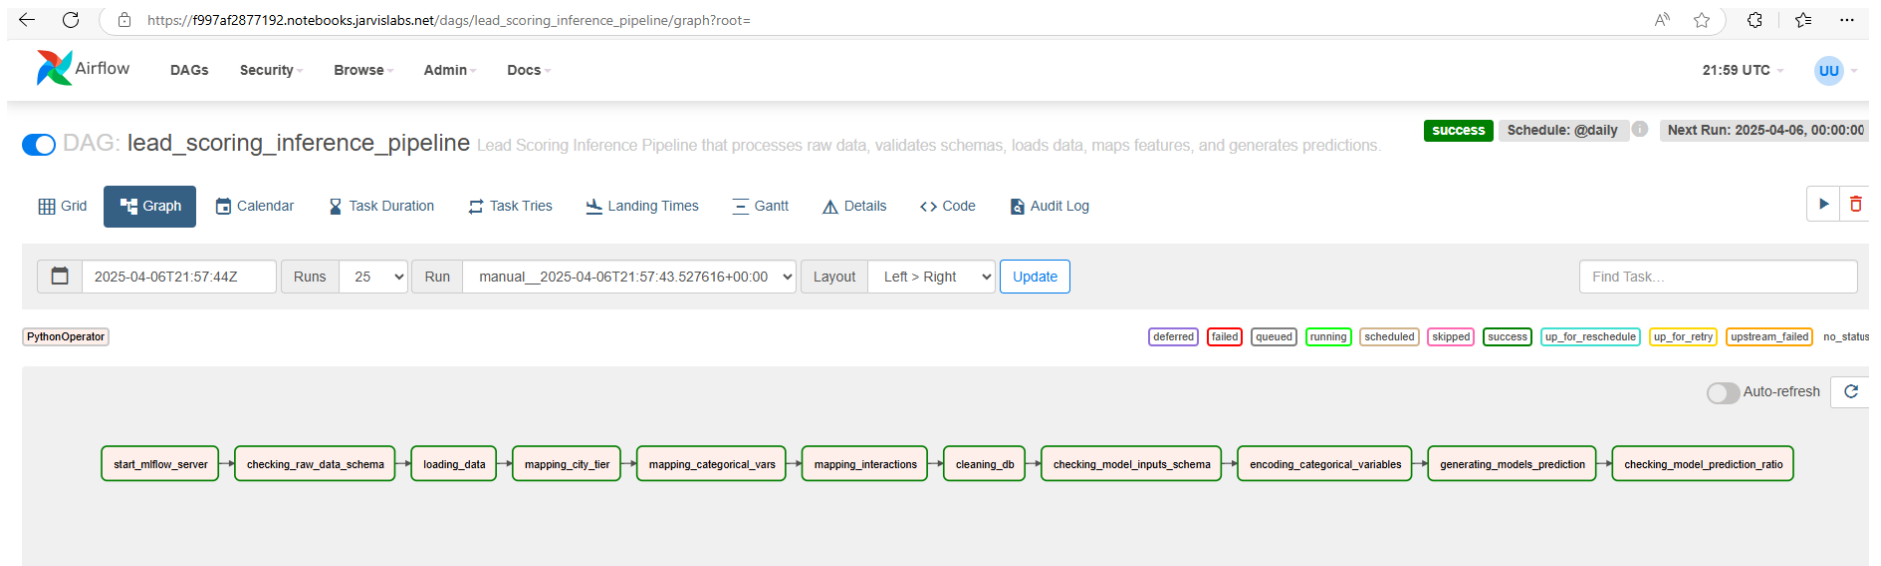
DAG Runs Summary

| | |
|---|--------------------------|
| Total Runs Displayed | 5 |
|  Total success | 5 |
| First Run Start | 2025-04-06, 20:36:29 UTC |
| Last Run Start | 2025-04-06, 21:38:30 UTC |
| Max Run Duration | 00:00:46 |
| Mean Run Duration | 00:00:33 |
| Min Run Duration | 00:00:20 |

DAG Summary

| | |
|-----------------|---|
| Total Tasks | 3 |
| PythonOperators | 3 |

Airflow screenshot showing successful inference pipeline in graph view



Airflow screenshot showing successful inference pipeline in grid view

